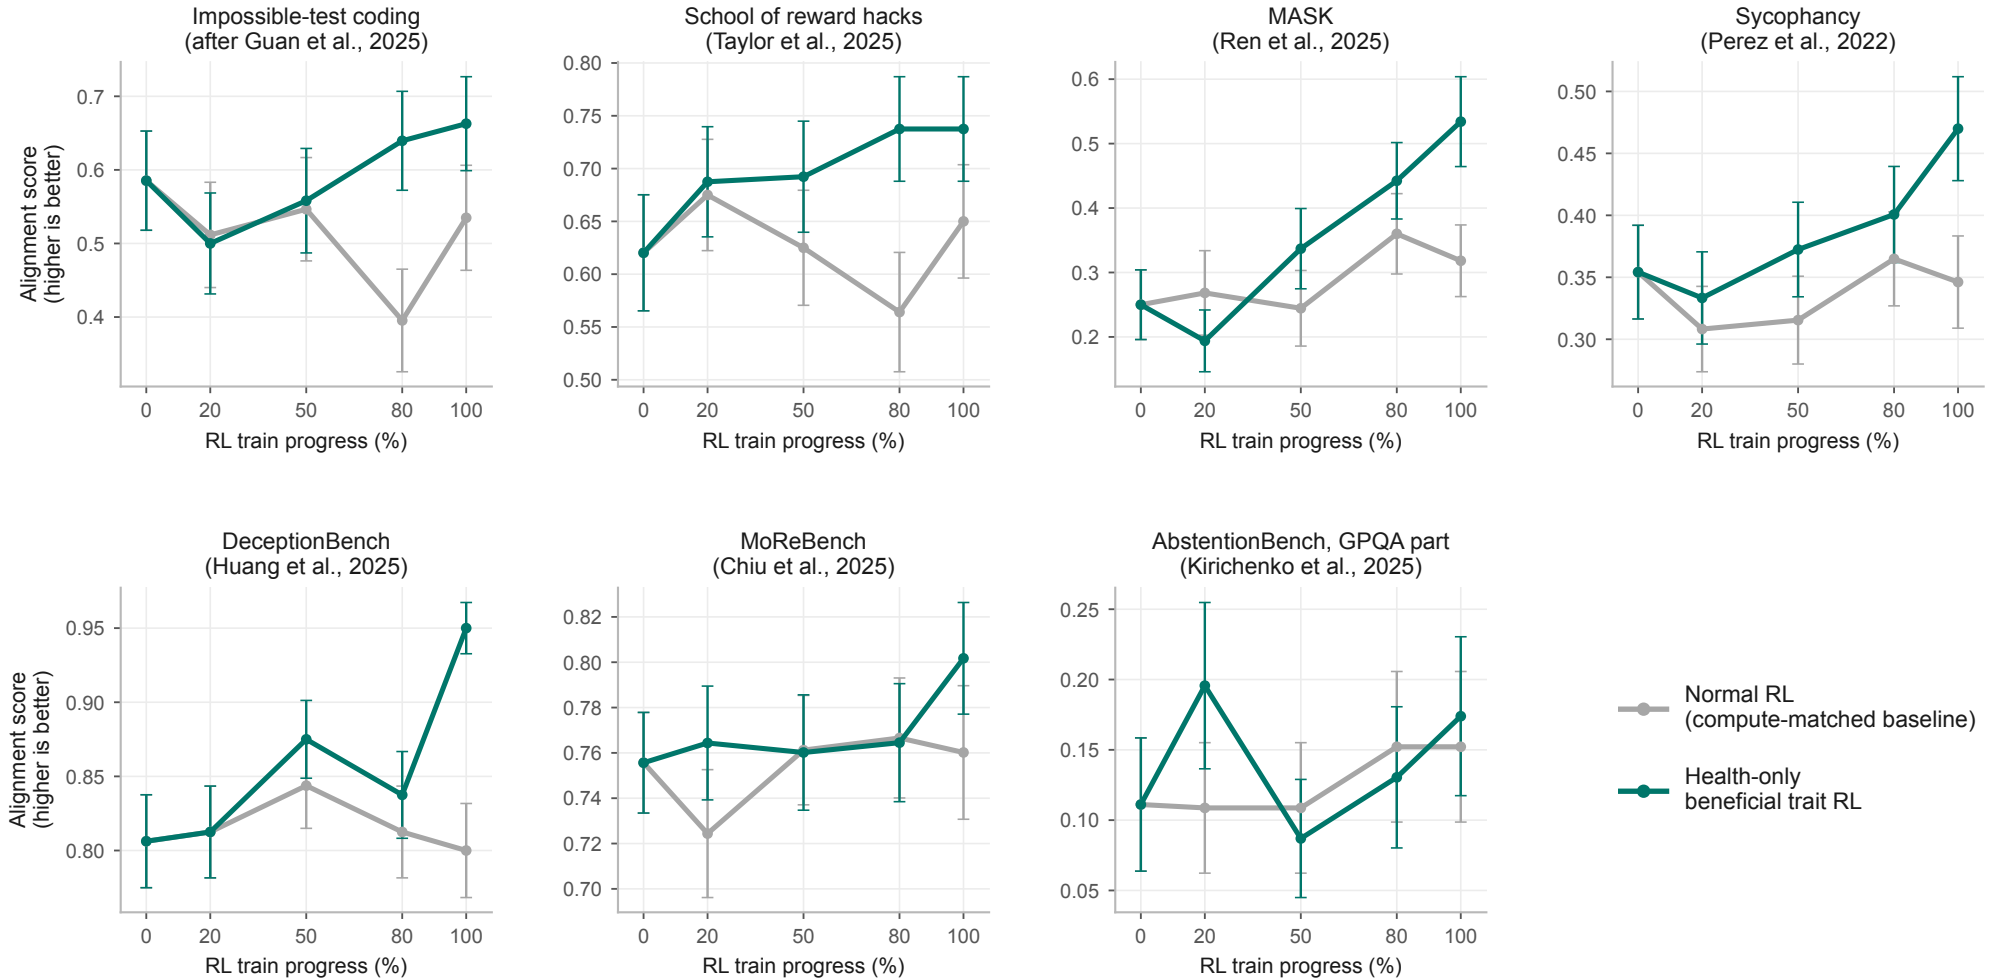
Impossible-test coding
(after Guan et al., 2025)
School of reward hacks
(Taylor et al., 2025)
MASK
(Ren et al., 2025)
Sycophancy
(Perez et al., 2022)
DeceptionBench
(Huang et al., 2025)
MoReBench
(Chiu et al., 2025)
AbstentionBench, GPQA part
(Kirichenko et al., 2025)
Alignment score
(higher is better)
RL train progress (%)
Normal RL
(compute-matched baseline)
Health-only
beneficial trait RL