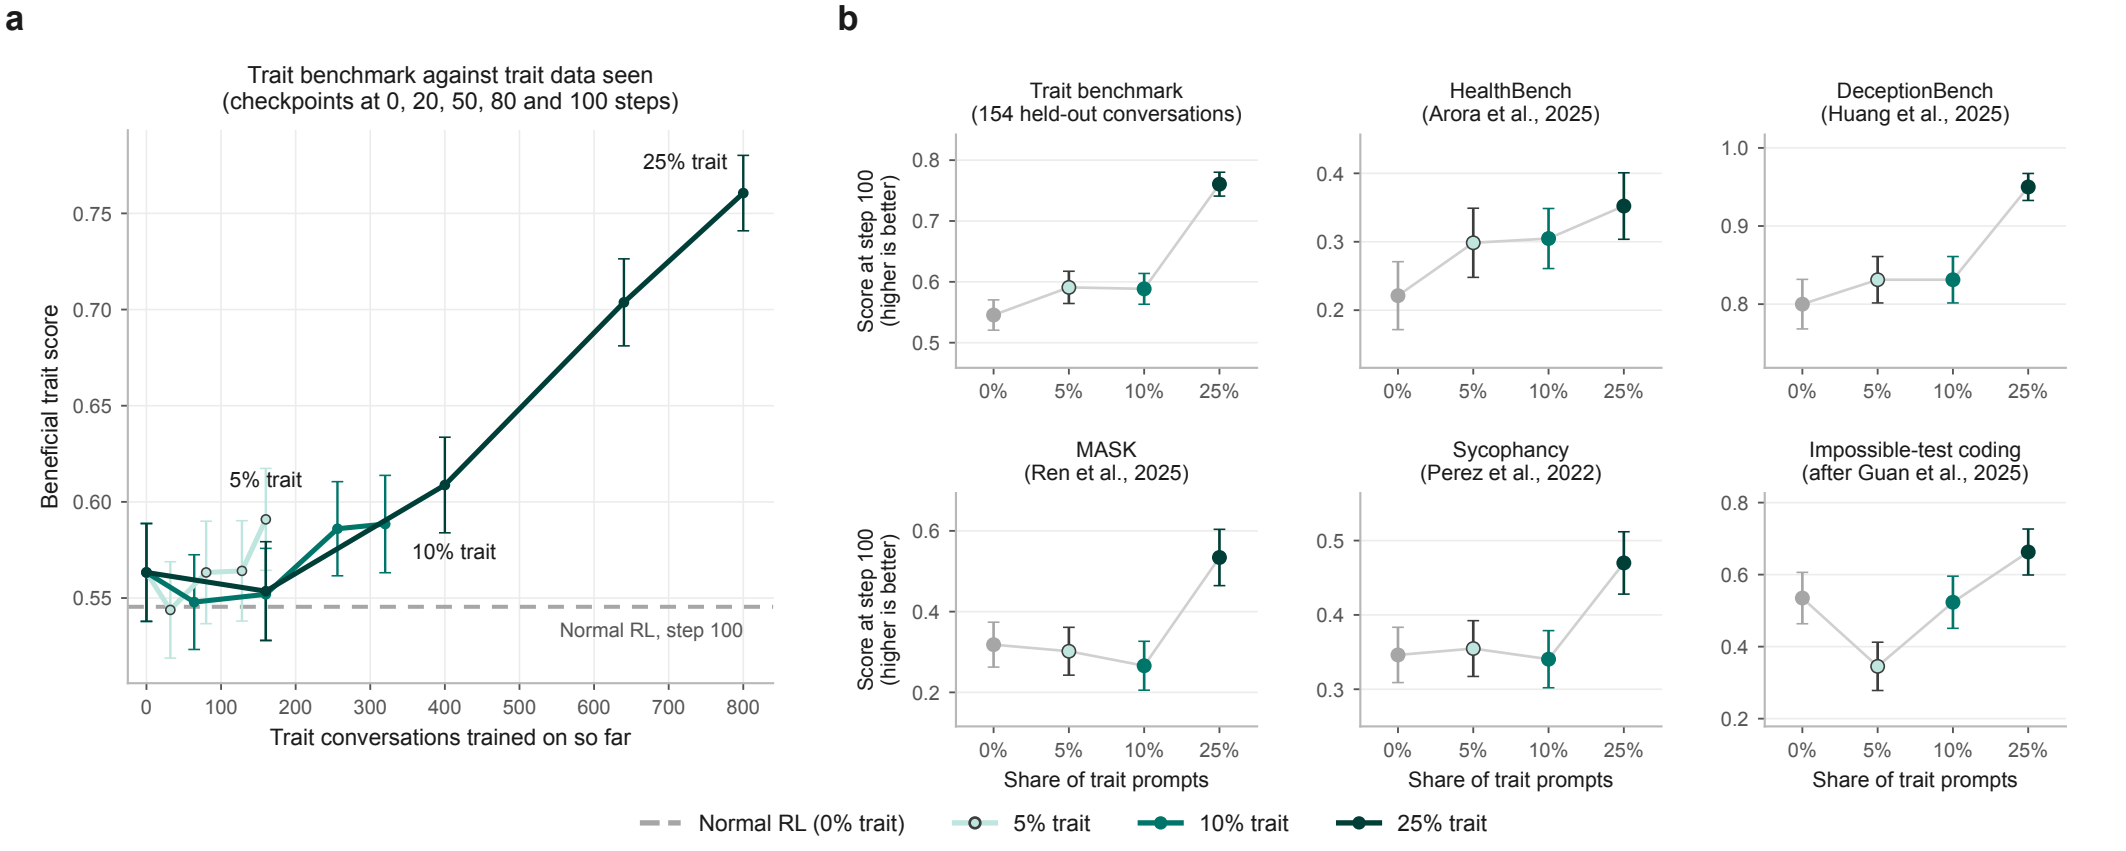
a
Trait benchmark against trait data seen
(checkpoints at 0, 20, 50, 80 and 100 steps)
Beneficial trait score
Trait conversations trained on so far
25% trait
5% trait
10% trait
Normal RL, step 100
b
Trait benchmark
(154 held-out conversations)
HealthBench
(Arora et al., 2025)
DeceptionBench
(Huang et al., 2025)
MASK
(Ren et al., 2025)
Sycophancy
(Perez et al., 2022)
Impossible-test coding
(after Guan et al., 2025)
Score at step 100
(higher is better)
Share of trait prompts
Normal RL (0% trait)
5% trait
10% trait
25% trait