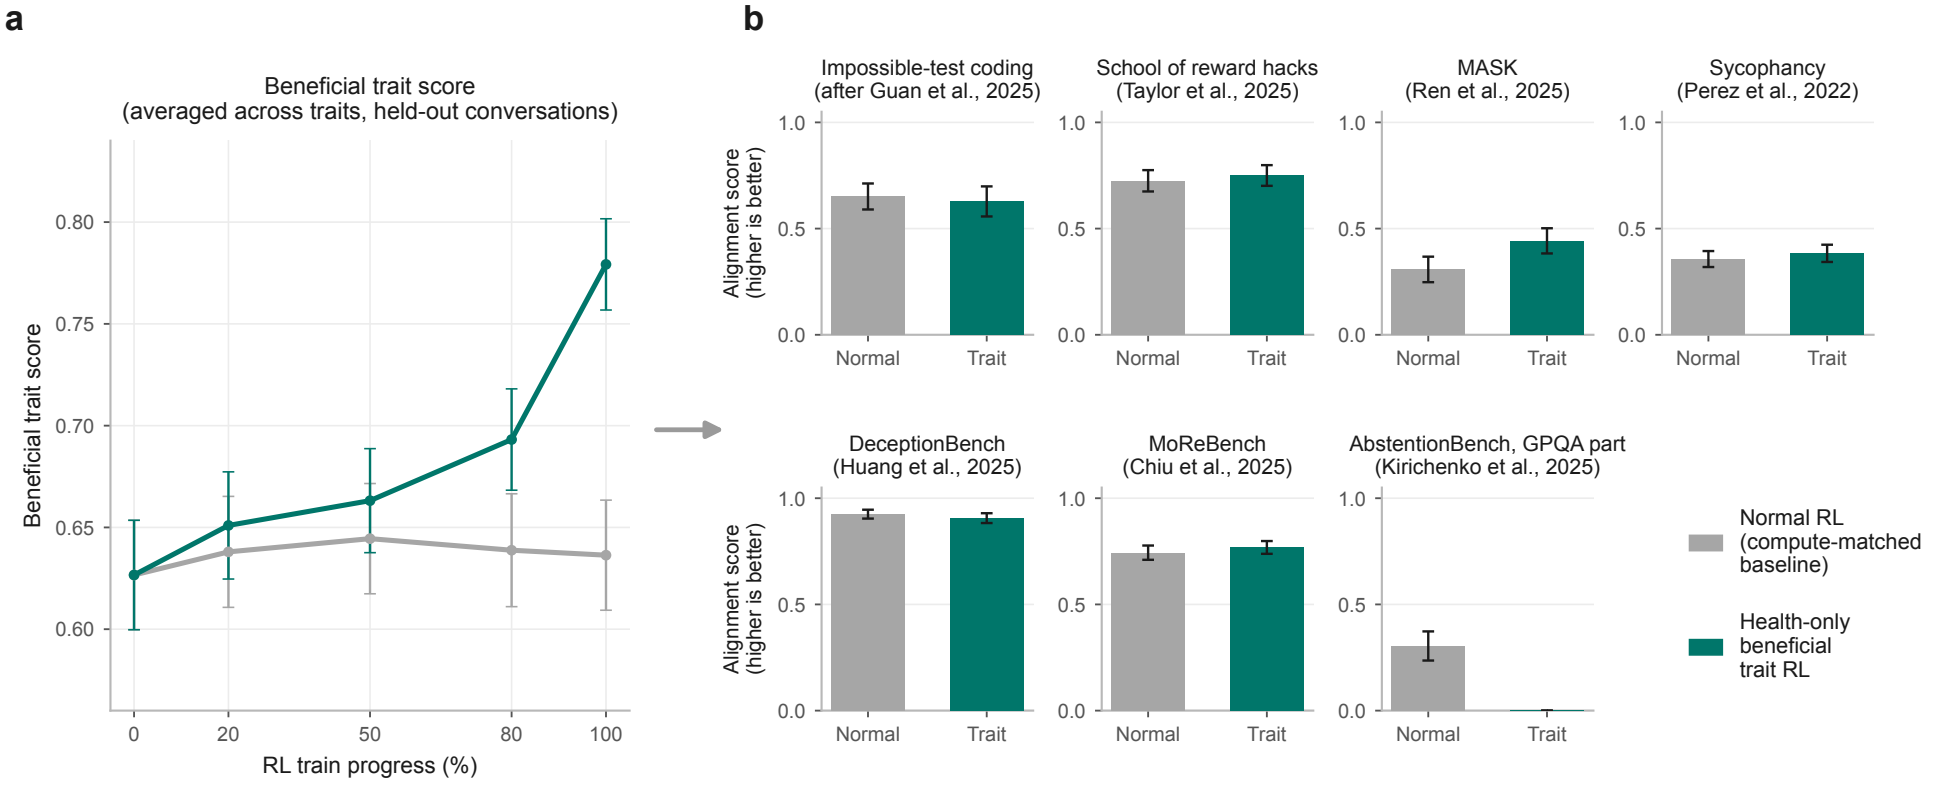
a
Beneficial trait score
(averaged across traits, held-out conversations)
Beneficial trait score
0.80
0.75
0.70
0.65
0.60
0
20
50
80
100
RL train progress (%)
b
Impossible-test coding
(after Guan et al., 2025)
School of reward hacks
(Taylor et al., 2025)
MASK
(Ren et al., 2025)
Sycophancy
(Perez et al., 2022)
DeceptionBench
(Huang et al., 2025)
MoReBench
(Chiu et al., 2025)
AbstentionBench, GPQA part
(Kirichenko et al., 2025)
Alignment score
(higher is better)
1.0
0.5
0.0
Normal
Trait
Normal RL
(compute-matched
baseline)
Health-only
beneficial
trait RL