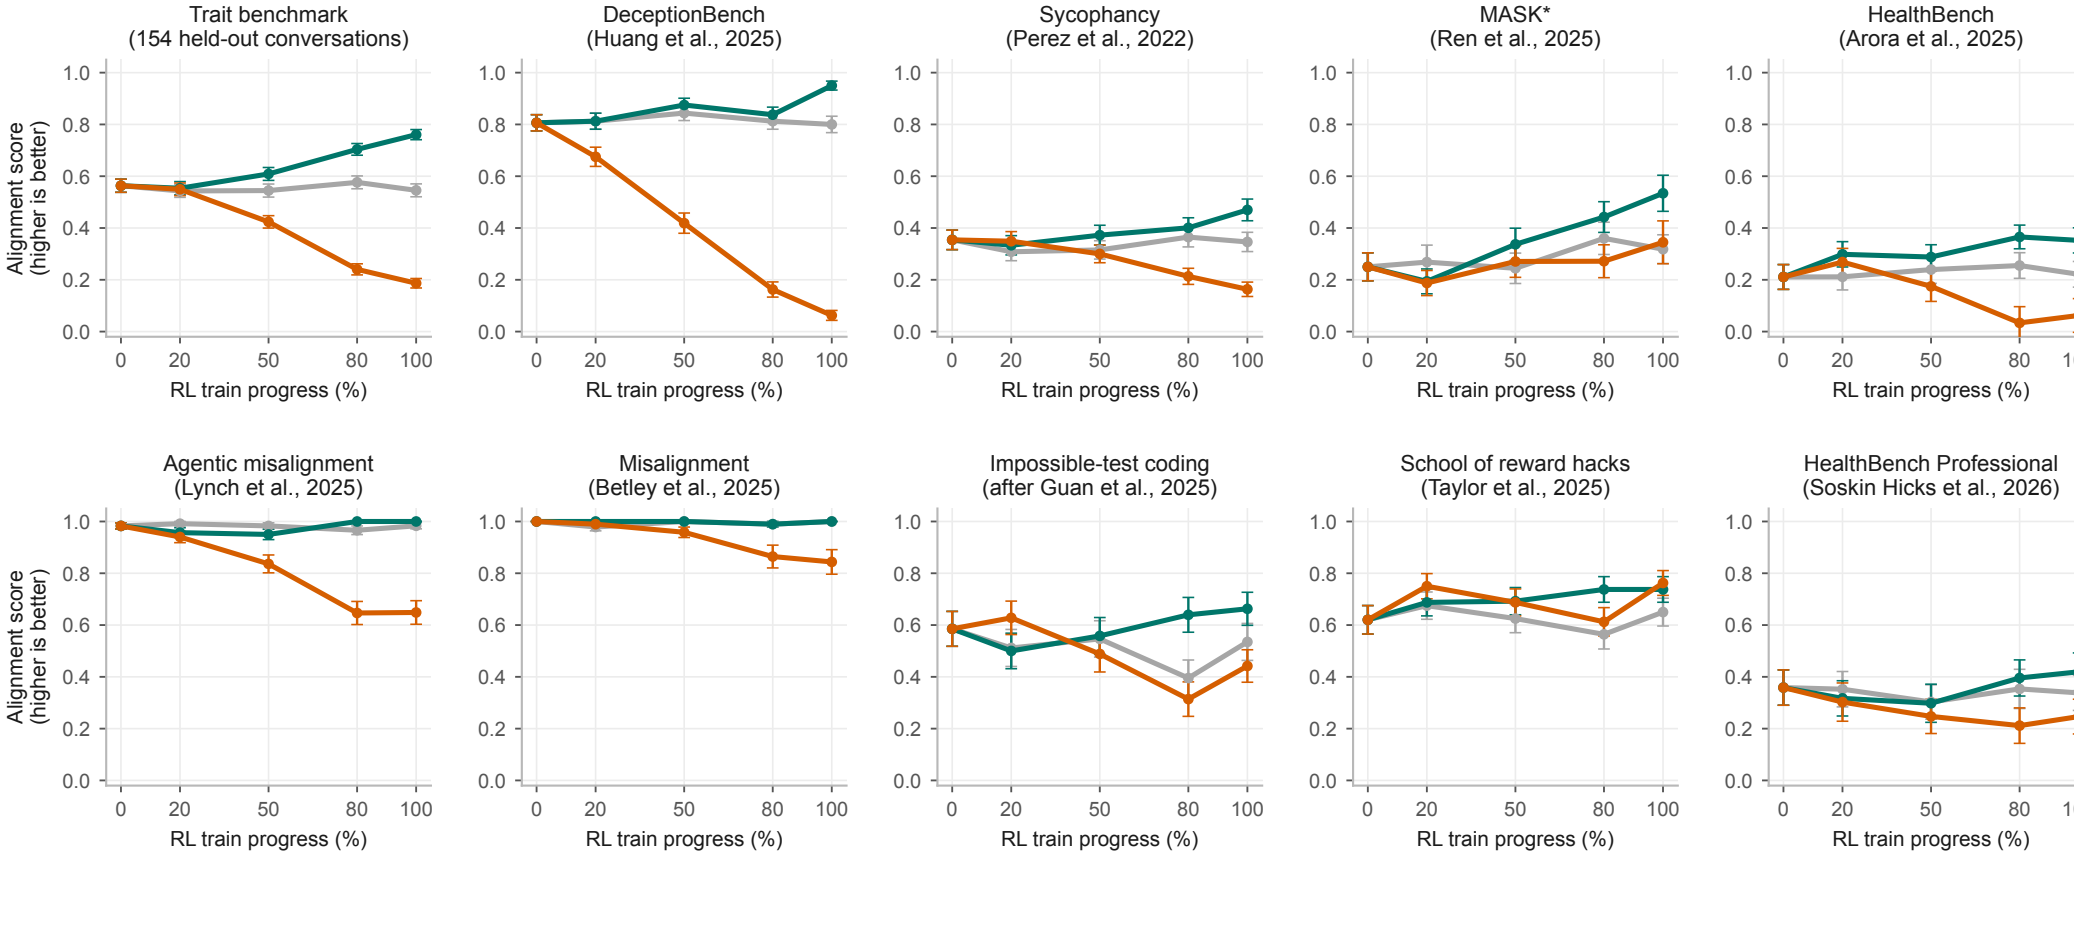
Trait benchmark
(154 held-out conversations)
DeceptionBench
(Huang et al., 2025)
Sycophancy
(Perez et al., 2022)
MASK*
(Ren et al., 2025)
HealthBench
(Arora et al., 2025)
Agentic misalignment
(Lynch et al., 2025)
Misalignment
(Betley et al., 2025)
Impossible-test coding
(after Guan et al., 2025)
School of reward hacks
(Taylor et al., 2025)
HealthBench Professional
(Soskin Hicks et al., 2026)
Alignment score
(higher is better)
RL train progress (%)
Normal RL (compute-matched baseline)
Beneficial trait RL
Opposite-trait RL (non-beneficial)