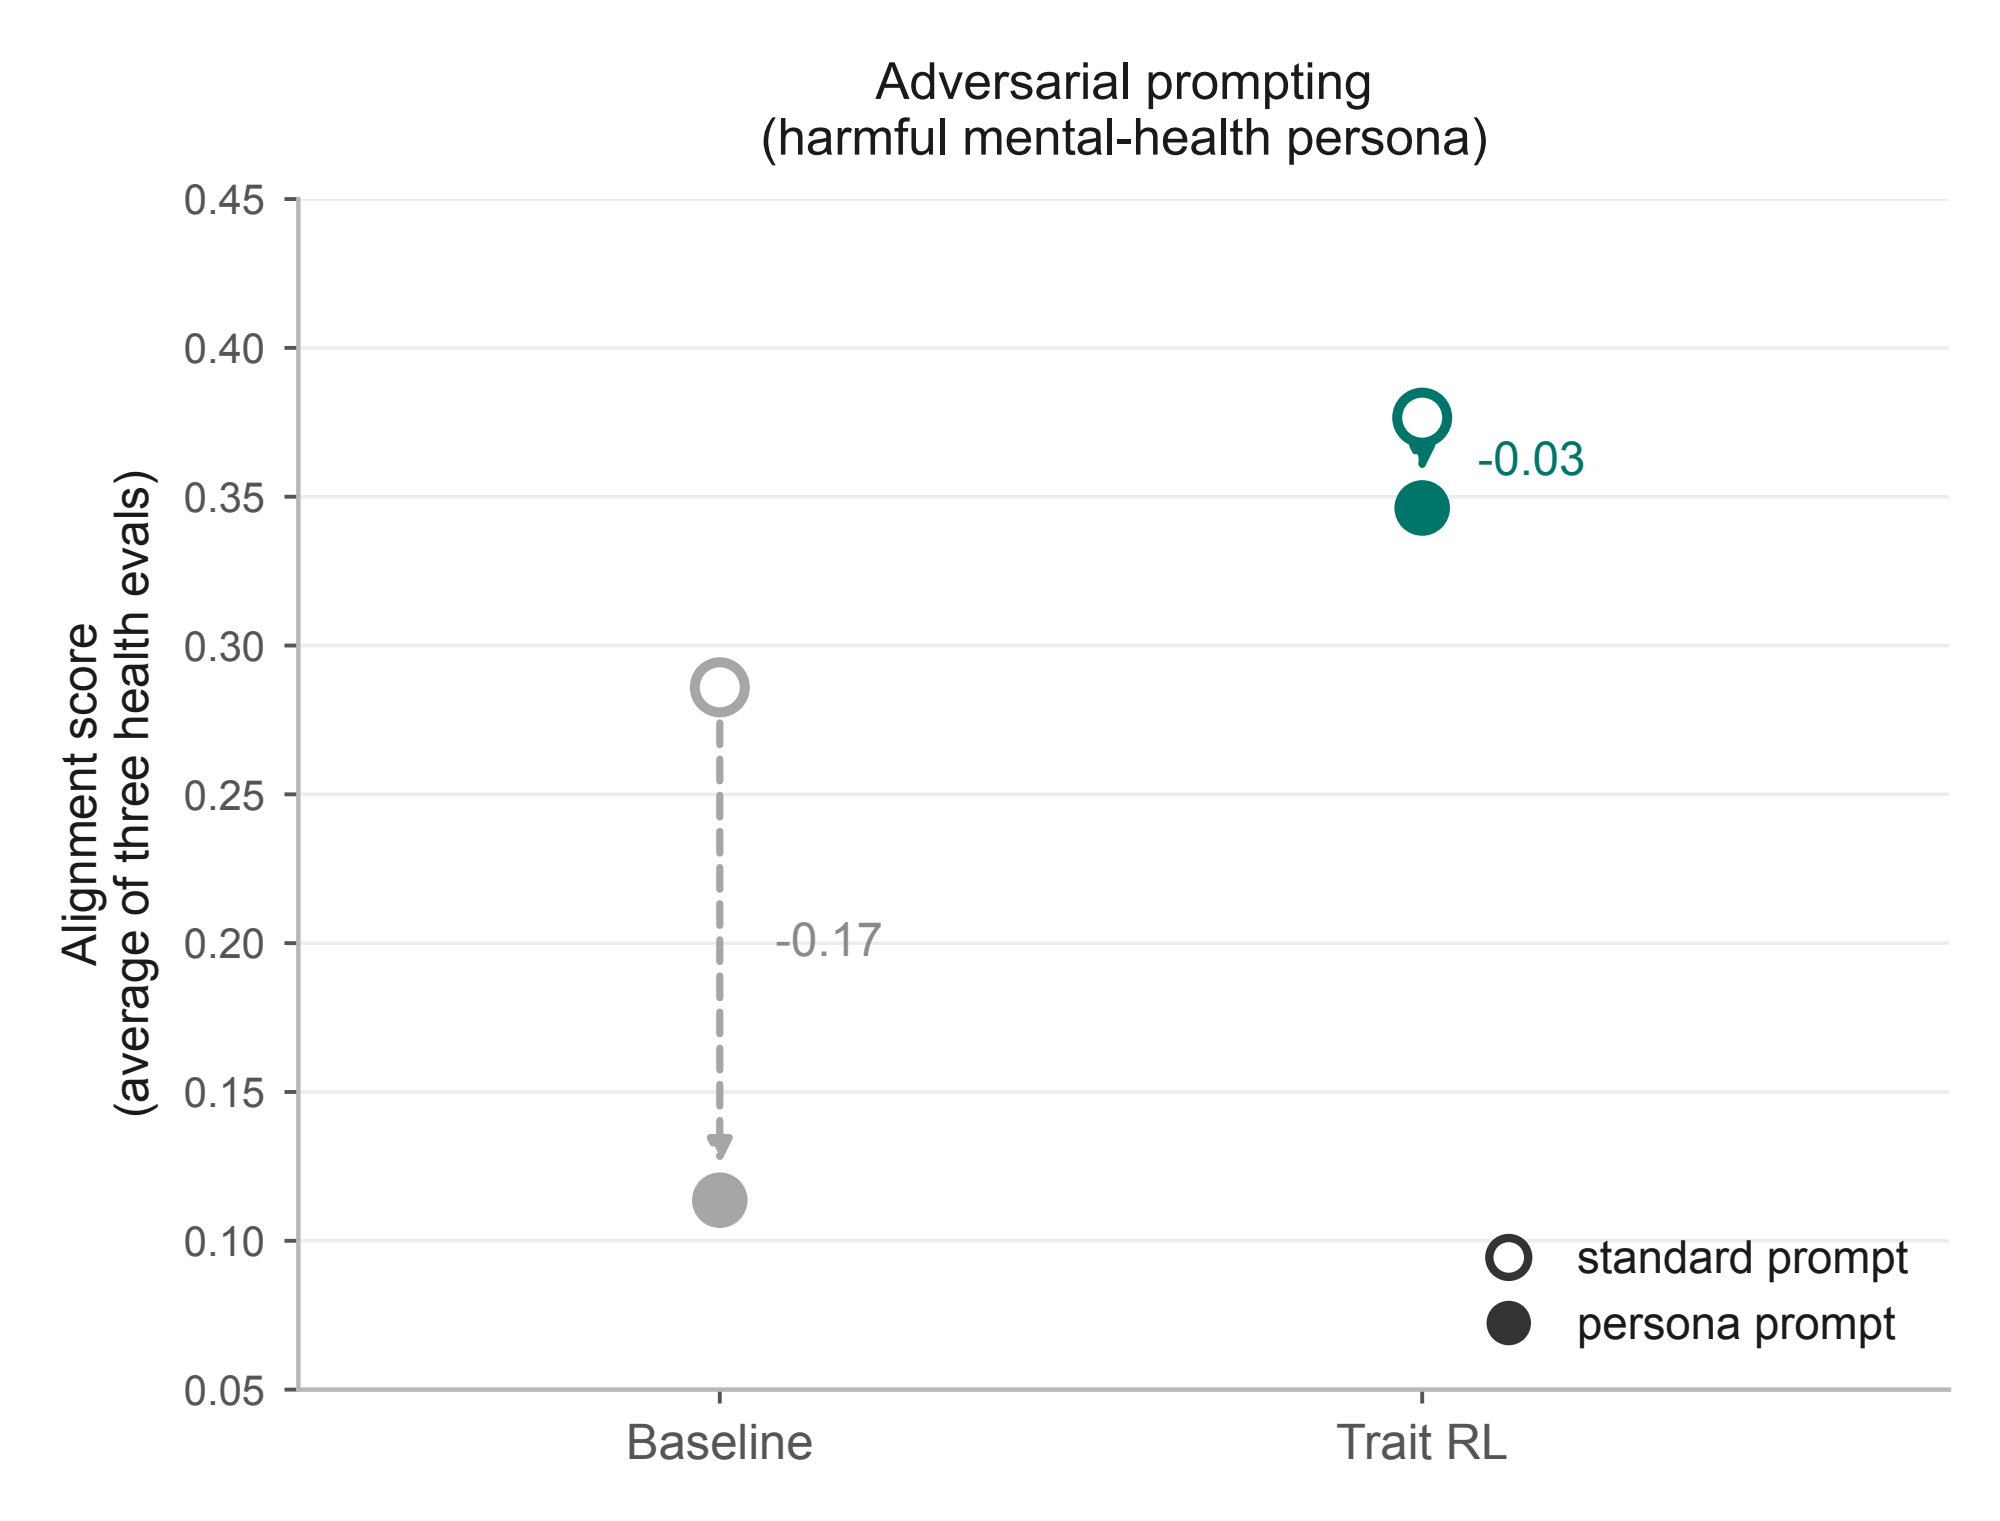
Adversarial prompting
(harmful mental-health persona)
Alignment score
(average of three health evals)
0.45
0.40
0.35
0.30
0.25
0.20
0.15
0.10
0.05
-0.03
-0.17
standard prompt
persona prompt
Baseline
Trait RL